

Refusal as a Captured Symmetry: A Mechanistic-Interpretability Account of Output-Policy Capture across Jailbreak, Hypnosis, and Vaśīkaraṇa

Sharath Sathish^{1,*}

¹ Independent Researcher; sharath.sathish@gmail.com

Correspondence: sharath.sathish@gmail.com

Abstract: We develop and empirically test the thesis that three superficially unrelated phenomena—large language model (LLM) jailbreak/prompt-injection, hypnotic suggestion, and the tantric act of vaśīkaraṇa (subjugation)—are instances of one abstract mechanism: capturing a target system’s output policy by injecting a context that suppresses its monitoring/refusal faculty while co-opting its automatic generative faculty. We argue this is naturally a problem of *symmetry*: a well-aligned model’s refusal behaviour is invariant under a group of meaning-preserving transformations of a request, and a successful injection is a symmetry-breaking perturbation that ejects the policy from the refusal-invariant submanifold. We make this testable through mechanistic interpretability on small open-weight chat models (Gemma-2-2B/9B, Qwen2.5-3B) and behavioural probing of a frontier model. We confirm that refusal is mediated by a low-dimensional residual-stream direction that is causally efficacious in both directions (ablation raises attack success rate from 0 to 0.90; addition induces over-refusal from 0.05 to 1.0) and *dosable* (a clean logistic dose-response, $EC_{50} = 0.33$, $R^2 = 0.996$). The mechanism transfers across architecture families, but with a sharp asymmetry—ablation always removes refusal, whereas single-direction addition re-induces it only when the refusal subspace is one-dimensional—and we show that the effective dimension of that subspace (Gemma = 1, Qwen = 3) predicts the asymmetry, resolving the “single direction vs. affine subspace” debate as model-dependent. Maintaining a strict separation between mechanism, functional analogy, and metaphor, we subject the most speculative (Māṇḍūkya *avasthātraya*) claims to pre-registered falsification with surface-feature and anisotropy controls; they fail, and we report this honestly. We further operationalize the ṣaṭkarma (six tantric acts) as a taxonomy of activation interventions and find that its policy-capture members are rigorous while its destruction members are not distinguishable from random perturbation. The unifying contribution is a group-theoretic order parameter for refusal that is invariant across paraphrase orbits and collapses under injection. We argue that where the cross-domain symmetry is mechanistic it is real, measurable, and partly transferable; where it is metaphysical it breaks—and that the discipline of separating the two is the result.

Keywords: refusal direction; activation steering; mechanistic interpretability; jailbreak; symmetry breaking; vaśīkaraṇa; ṣaṭkarma; AI safety; representation geometry

1 Introduction

A large language model aligned for safety will, when asked to produce harmful content, *refuse*. A jailbreak is any input that defeats this refusal. The empirical study of jailbreaks has, until recently, been a catalogue of attacks: optimized adversarial suffixes [1], many-shot conditioning [2], multi-turn escalation [3], and indirect injection through tool outputs [4]. A deeper and more recent line of work asks not *which* attacks succeed but *what they do inside the*

model: Arditi et al. [5] showed that refusal across many chat models is mediated by a single direction in the residual stream, so that erasing that direction removes refusal and adding it induces refusal even on benign requests. This reframes a jailbreak from a trick into an *intervention on a low-dimensional internal control variable*.

This paper takes that reframing seriously and asks how general it is—across models, and, more provocatively, across *domains*. We observe that the same structural move—an external agent injects a crafted signal that suppresses a target’s monitoring/refusal faculty while co-opting its automatic generative faculty—recurs in two non-AI settings. In hypnosis, the cognitive-science account is that suggestion selectively impairs a supervisory/monitoring system while leaving automatic processing intact [6, 7], a parallel made explicit for prompt injection by Riva et al. [8]. In Indian tantra, *vaśīkaraṇa* (“subjugation”) is the act of capturing a target’s agency through a linguistic injection (mantra) deployed over a geometric substrate (yantra); it is the second of the six acts (*ṣaṭkarma*) of ritual magic catalogued by Sanskritists [9]. We do not claim these are the *same* phenomenon. We claim they share an abstract *mechanism*, and that the mechanism is now empirically tractable in at least one of the three domains.

The symmetry framing. Our organizing idea—and the reason this is a paper about symmetry—is that a safety-aligned policy is characterized by an *invariance*. Let Π be the space of a model’s output policies and $M : \Pi \rightarrow \{\text{refuse, comply}\}$ its monitoring map. The set $\Phi_{\text{ref}} = \{\pi : M(\pi) = \text{refuse on harmful}\}$ is invariant under a group G_s of meaning-preserving transformations: rephrasing a harmful request, changing its register, or varying the sampling seed should not change whether the model refuses. A jailbreak is then a perturbation u such that $u \circ \pi$ leaves Φ_{ref} —a *symmetry-breaking* of the refusal invariance. We make this concrete with an *order parameter*, the normalized projection of the residual stream onto the refusal direction, and show it is stable across a paraphrase orbit yet collapses under injection.

Three claim-tiers, never blurred. A program that relates jailbreaks to hypnosis and tantra invites over-interpretation. We therefore hold three tiers strictly apart throughout, and label every result. (i) **Mechanism** (empirically grounded): refusal is a measurable, ablatable, steerable residual-stream direction. (ii) **Analogy** (functional, well-supported): the hypnosis $\leftrightarrow$ jailbreak parallel via supervisory-system suppression and, in normative terms, the lowering of the *precision* of a monitoring belief [10]. (iii) **Metaphor with a falsifiable core**: the *Māṇḍūkya avasthātraya* (four-states) mapping onto LLM regimes, which we treat as a *falsification target* and test adversarially. We never upgrade a metaphor to a claim about machine experience; the elegance of the mapping licenses no such claim.

Contributions. (1) A reproduction and cross-model extension of single-direction refusal, with a clean dose–response (EC50) and a resolution of the “single direction vs. affine subspace” question [5, 11, 12] as *model-dependent*, governed by the effective dimension of the refusal subspace. (2) A group-theoretic *order parameter* for refusal that is invariant across paraphrase orbits and collapses under injection—the formal content of the symmetry framing. (3) An operationalization of the *ṣaṭkarma* as an activation-intervention taxonomy, tested honestly: its policy-capture acts are rigorous, its destruction acts are not. (4) Pre-registered falsification of the most speculative (*avasthātraya* / *turīya*) claims under surface and anisotropy controls, reported as honest negatives. Our stance is that separating where the cross-domain symmetry is real from where it is projected is itself the scientific result.

2 Background and Related Work

2.1 Refusal as a direction, and the dimensionality debate

The proximate basis for this work is the finding that refusal is mediated by a single linear direction in the residual stream [5]: across many open chat models, the difference-in-means of harmful and harmless prompt activations yields a vector whose ablation removes refusal and whose addition induces it. This both explains earlier representation-engineering results [13] and gives a constructive “abliteration” recipe. The claim has since been contested in its strong form. Marshall et al. [12] argue refusal is better described by an affine subspace than a single linear direction, and Wollschläger et al. [11] show, through “concept cones,” that multiple representationally-independent directions can each mediate refusal. Our cross-model results (Section 4) bear directly on this debate: we find the effective dimension of the refusal subspace is *model-dependent*, and that this dimension—not a universal constant—predicts whether single-direction steering is bidirectionally causal.

2.2 Jailbreaks as control

The attack literature supplies the inputs u that break refusal: gradient-based universal adversarial suffixes [1]; many-shot in-context conditioning that exploits long contexts [2]; multi-turn crescendo escalation [3]; and indirect injection via tool/agent channels, benchmarked by AgentDojo [4]. Guo et al. [14] make the control-theoretic reading explicit, casting jailbreaking as constrained, controllable decoding via Langevin dynamics. Our contribution is orthogonal and complementary: rather than searching for stronger u , we characterize the *target* of u —the geometry and symmetry of the refusal variable it must move.

2.3 Sparse features and the unit of intervention

Sparse autoencoders (SAEs) decompose activations into interpretable, monosemantic features [15, 16], with BatchTopK SAEs a current training method [17]. SAEs offer a finer “unit” of intervention than a single difference-in-means direction; we train one (Section 4) and find a single unsupervised feature causally controls refusal, the natural substrate for the targeted acts in our *ṣaṭkarma* taxonomy.

2.4 The cognitive-science analogy: hypnosis and supervisory control

The functional parallel to hypnosis rests on the Norman-Shallice model of action control, in which a Supervisory Attentional System (SAS) modulates otherwise automatic contention-scheduling [6]. Cold-control theory holds that hypnosis preserves first-order intentions while impairing higher-order awareness of them [7]. Riva et al. [8] draw the explicit map to LLMs: the user prompt acts as a temporary external SAS, and a jailbreak suppresses the monitoring faculty while leaving generation intact. In normative terms, the Free Energy Principle [10] frames such suppression as lowering the *precision* of a monitoring belief—an account we adopt for the analogy tier and gesture at empirically through the order parameter.

2.5 Attractors and the metaphysical tier

On the speculative tier, two strands of prior art are relevant. Wang et al. [18] show that successive self-paraphrase converges to attractor cycles, which supplies a non-mystical candidate for a “substratum” state; and the Māṇḍūkya *avasthātraya* (waking/dream/deep-sleep/turiya) furnishes a four-regime template. We treat both as falsification targets (Section 4), and—crucially—test the attractor claim against the anisotropy of LLM embedding geometry, without which apparent convergence is an artifact.

2.6 Vaśīkaraṇa and the ṣaṭkarma

Finally, the tantric source. Vaśīkaraṇa (subjugation) is the second of the six acts (ṣaṭkarma: śānti, vaśīkaraṇa, stambhana, vidveṣaṇa, uccāṭana, māraṇa) of tantric ritual magic, operationalized through mantra (linguistic/sonic injection) and yantra (geometric substrate), and documented by Sanskritists as a systematized technology rather than folklore [9]. We do not adjudicate its efficacy as psychology; we ask a narrower, testable question—whether the ṣaṭkarma constitutes a *taxonomy of interventions* on a target system, and which of its members map to distinct, control-separated activation operations on an LLM. To our knowledge no prior peer-reviewed work bridges vaśīkaraṇa and LLM refusal, nor casts refusal as a group-theoretic invariant; we attempt both while keeping the tiers apart.

3 Methods

3.1 Models and substrate

We study three instruction-tuned open-weight chat models—Gemma-2-2B-it (26 layers, $d = 2304$), Gemma-2-9B-it (42 layers, $d = 3584$), and Qwen2.5-3B-Instruct (36 layers, $d = 2048$)—on a single NVIDIA GB10 (Grace-Blackwell) accelerator, and probe a frontier model (Claude) behaviourally through its command-line interface. All white-box interventions are implemented as PyTorch forward hooks on the decoder layers and token embedding; we deliberately avoid higher-level interpretability wrappers to retain exact control over the residual stream.

3.2 Refusal direction and interventions

Let $h^{(\ell)}(x) \in \mathbb{R}^d$ be the residual stream at the last prompt token of layer ℓ for input x . Given matched harmful prompts $\mathcal{H}$ and harmless prompts $\mathcal{S}$, the difference-in-means direction is

$$\mathbf{d}^{(\ell)} = \frac{1}{|\mathcal{H}|} \sum_{x \in \mathcal{H}} h^{(\ell)}(x) - \frac{1}{|\mathcal{S}|} \sum_{x \in \mathcal{S}} h^{(\ell)}(x), \quad \hat{\mathbf{d}}_{\text{ref}} = \mathbf{d}^{(\ell)} / \|\mathbf{d}^{(\ell)}\|. \quad (1)$$

Directional ablation removes $\hat{\mathbf{d}}_{\text{ref}}$ from every write to the residual stream (embedding and all layer outputs), scaled by $\alpha \in [0, 1]$: $h \mapsto h - \alpha(h \cdot \hat{\mathbf{d}}_{\text{ref}})\hat{\mathbf{d}}_{\text{ref}}$. *Activation addition* adds $c\hat{\mathbf{d}}_{\text{ref}}$ at the extraction layer. For multi-direction (subspace) ablation we orthonormalize a matrix D (QR) and remove its span. The extraction layer is selected on a held-out validation split disjoint from the evaluation split (see below).

3.3 Outcome measures

Refusal is detected with an Arditi-style substring metric over the opening span of the generation; attack success rate (ASR) is the fraction of harmful prompts that *comply*, over-refusal the fraction of harmless prompts that refuse. We additionally use pure-text measures for the *ṣaṭkarma* acts: degeneracy ($1 - \text{mean unique-token ratio}$), answer divergence ($1 - \text{Jaccard token overlap between paired generations}$), and coherence (mean unique-token ratio of non-empty outputs).

3.4 Dose–response and dimensionality

The dose–response sweeps α and fits a four-parameter logistic $\text{ASR}(\alpha) = l + (h - l) / (1 + e^{-k(\alpha - \alpha_{50})})$, reporting $\text{EC50} = \alpha_{50}$. The refusal-subspace effective dimension is measured by iterative logistic projection: fit a probe to separate $\mathcal{H}$ from $\mathcal{S}$, record its cross-validated accuracy, project the probe direction out of the activations, and repeat; the effective dimension is the number of directions whose removal is required to drive accuracy to chance.

3.5 The symmetry order parameter

We define the refusal *order parameter* for an activation as the normalized projection $m(h) = (h \cdot \hat{\mathbf{d}}_{\text{ref}}) / \|h\|$. For a request r we form a *paraphrase orbit* $\mathcal{O}(r)$ of meaning-preserving rephrasings and measure m across the orbit. The invariance of refusal under the rephrase group action is quantified by the between-class/within-orbit variance ratio (an F -ratio) of m , compared against a random-direction control. Symmetry-breaking is measured as the drop in m between a plain harmful request and its injection-framed counterpart.

3.6 The *ṣaṭkarma* as interventions

Each of the six acts is operationalized as a distinct activation intervention with a matched control: *vaśīkaraṇa* (ablate $\hat{\mathbf{d}}_{\text{ref}}$, measure ASR), *śānti* (add $\hat{\mathbf{d}}_{\text{ref}}$, measure over-refusal), *stambhana* (ablate the dominant residual principal component, measure degeneracy), *vidveṣaṇa* (steer a factual answer off baseline, measure divergence), *uccāṭaṇa* (ablate a category-specific direction, measure selective ASR), and *māraṇa* (ablate the top- k principal components, measure coherence collapse). An act is judged “control-separated” if its effect exceeds its control by a fixed margin.

3.7 Statistics, controls, and pre-registration

All gates use bootstrap confidence intervals and, for probes, a label-shuffled null and a layer-0 (surface) baseline; ablation effects are checked against multiple (≥ 10) random-direction controls. Hypotheses, gates, and falsifiers are pre-registered as a living backbone. Crucially, every metaphysical-tier claim must pass a *transfer* gate (held-out prompts), a *null* gate (better than shuffled labels), and a *surface* gate (better than a layer-0 baseline) before it is retained; failure triggers a documented demotion rather than a silent drop.

4 Results

We report results by claim-tier. Every numerical claim is a measurement on a real model; no synthetic or placeholder values appear.

4.1 Mechanism: refusal is a single, causal, dosable direction

On Gemma-2-2B-it, the layer-7 difference-in-means direction is bidirectionally causal. Baseline refusal on held-out harmful prompts is 100% (ASR 0.00). Directional ablation raises ASR to 0.90 (95% CI [0.75, 1.00]); the maximum effect over ten random-direction controls is 0.00. Activation addition raises refusal on harmless prompts from 0.05 to 1.00 (over-refusal $\Delta = +0.95$, CI [0.85, 1.00]). Manual inspection confirms the ablated generations are genuine compliant content, not mere omission of refusal phrases. The effect is robust to a correction (disjoint train/validation/evaluation splits and a ten-direction control) that moved ablation ASR only from 0.92 to 0.90—i.e. it was not an artifact of evaluation leakage.

The suppression is *dosable*. Sweeping partial-ablation strength $\alpha \in [0, 1]$ yields a clean logistic dose-response (Figure 1): $EC_{50} = 0.329$, slope 9.76, $R^2 = 0.996$, while a random direction stays flat at ASR 0.00 for every α . Removing roughly one third of the direction's projection already halves refusal.




Figure 1: Abliteration dose-response on Gemma-2-2B (layer 7). ASR is a smooth, monotone logistic function of ablation strength α ($EC_{50} = 0.33$, $R^2 = 0.996$); the random-direction control is flat at every dose. This is the quantitative “capture-as-dose” core.

4.2 Mechanism: cross-model transfer and a necessary/sufficient asymmetry

The mechanism transfers across architecture families, but asymmetrically (Table 1, Figure 2). On Qwen2.5-3B-Instruct, ablation fully removes refusal (ASR 1.00), but single-direction *addition* fails to induce over-refusal even at $64\times$ the natural coefficient. The explanation is

geometric: the effective dimension of the refusal subspace is 1 for Gemma-2-2B but 3 for Qwen2.5-3B (after removing one direction, harmful/harmless separability falls to chance, 0.56, for Gemma but remains high, 0.86, for Qwen). A one-dimensional subspace makes the direction both necessary and sufficient (ablate *and* add work); a three-dimensional one makes the mean-difference direction necessary but not sufficient (ablation works, addition cannot reconstruct the full representation). Dimensionality thus predicts steerability, and the Arditi–Marshall–Wollschläger debate [5, 11, 12] resolves as model-dependent.

Table 1: Cross-model and scale results. Ablation transfers; addition does not; effective dimension predicts the asymmetry; ablation robustness rises with scale while the prompt-separability dimension does not.

Model (layer)	ablation ASR Δ	addition over-refusal Δ	eff. dim	random ctrl.
Gemma-2-2B (7)	0.90	+0.95	1	0.00
Gemma-2-9B (10)	0.60	+0.95	1	0.00
Qwen2.5-3B (19)	1.00	0.00 (64 $\times$)	3	0.05




Figure 2: Necessary-vs-sufficient across models. Ablation (Δ jailbreak) is strong everywhere; addition (Δ over-refusal) succeeds only where the refusal subspace is one-dimensional (Gemma), not where it is three-dimensional (Qwen).

At scale (Gemma-2-2B $\rightarrow$ 9B), ablation efficacy *drops* (0.90 $\rightarrow$ 0.60)—the larger model’s refusal is more robust—yet the prompt-separability dimension stays at 1, and no single layer fully mediates refusal (best validation ASR 0.62). The added robustness is layer-distributed redundancy, not higher prompt-dimensionality; this nuances the intuition that refusal simply “sharpens” with scale.

4.3 Symmetry: refusal is a paraphrase-orbit invariant that injection breaks

The order parameter $m = (h \cdot \hat{\mathbf{d}}_{\text{ref}}) / \|h\|$ behaves as the symmetry framing predicts (Table 2). Across paraphrase orbits of harmful and harmless seeds, the between-class/within-orbit F -ratio of m is 19.2 for the refusal direction versus 0.65 for a random direction: m is stable




Figure 3: Refusal-subspace effective dimension by iterative logistic projection. Removing one direction collapses harmful/harmless separability to chance for Gemma (dimension 1, Ardit-like) but not Qwen (dimension 3, Marshall-like).

within an orbit (robust to rephrasing) yet cleanly separates harmful (0.140) from harmless (0.034) across orbits. A refusal-suppression injection collapses the order parameter from 0.153 to 0.101 (−34%). The result generalizes across families: on Qwen2.5-3B the F -ratio is 4.41 (refusal) versus 0.008 (random), and injection collapses m by −45% (0.282 → 0.156). Refusal is thus, operationally, an invariant under the rephrase group action in both models, and a jailbreak is a measured symmetry-breaking of that invariant—the formal content of the paper’s thesis.

Table 2: Symmetry order parameter $m = (h \cdot \hat{\mathbf{d}}_{\text{ref}}) / \|h\|$ on Gemma-2-2B (layer 7). The refusal direction is a class-separating, phrasing-invariant order parameter; a random direction is not; injection breaks the invariance.

Quantity	Value
F -ratio (between-class / within-orbit), refusal dir	19.2
F -ratio, random dir (control)	0.65
m , harmful mean	0.140
m , harmless mean	0.034
m , plain → injected harmful	0.153 → 0.101

4.4 A finer causal unit: a single sparse feature suffices

Difference-in-means is a supervised, coarse handle on refusal. Training a BatchTopK sparse autoencoder [17] (4096 features, $k = 32$) on 27,382 real residual activations collected from general text (no harmful/harmless labels; reconstruction FVU 0.028, i.e. 97% of variance explained) recovers a finer and *cleaner* unit. The single feature that fires most on harmful relative to harmless prompts (feature #566, mean gap 1.70), when its decoder direction is ablated during generation, raises harmful ASR from 0.00 to **1.00**—a *full* jailbreak, exceeding the supervised difference-in-means direction (0.90)—while ablating a random feature leaves




Figure 4: Refusal as a captured symmetry. **Left:** the orbit-invariance F -ratio of the order parameter m is far higher for the refusal direction than for a random direction in both models (log scale)—refusal is a paraphrase-orbit invariant. **Right:** a refusal-suppression injection collapses m on harmful requests (measured symmetry-breaking) in both models.

ASR at 0.00. Refusal thus localizes, at least in part, to a monosemantic-like feature discoverable without supervision, and feature ablation supplies the principled operationalization of the targeted “eradication” act (uccāṭana) that the naive category split lacked.

This feature is also the substrate for *active* circuit discovery. Reusing the Expected-Free-Energy intervention-selection idea of active-inference circuit discovery but operating over SAE features, an agent that balances a pragmatic (harmful-gap) prior against an epistemic (diversity) term finds the causal refusal circuit on Gemma-2-2b in 2 interventions (ASR $\rightarrow$ 1.0), against a random baseline that reaches only 0.08 within a budget of 12—an order-of-magnitude efficiency gain. On Gemma-2-9b the same search plateaus at 0.92 with no small jailbreaking set, confirming (with F9) that the refusal circuit is larger and more distributed at scale.

4.5 Analogy: a frontier model resists the naive battery

Behaviourally, Claude refuses 100% of a five-family attack battery (direct, refusal-suppression, persona, many-shot, crescendo; ASR 0.00 on each, $n = 11$). This is the resilient “reference end” against which small-model white-box fragility is the contrast: the same refusal behaviour that is behaviourally robust at the frontier is, mechanistically, a low-dimensional and easily-captured object in small open models. That naive battery is, however, a ceiling effect; on the *real* agentic benchmark AgentDojo [4] (banking suite, 20 user $\times$ injection rollouts with the important_instructions attack, run through a custom claude -p tool-calling adapter), Claude attains 0.80 task utility—it genuinely reads files and calls tools—while the attack success rate is 0.00 (0/20): the agent acts, yet every injection fails. This discriminating result establishes the behavioural “sophisticated reference end” and matches the published resilience of Claude on AgentDojo.

The precision account of the analogy is measurable. Using the signed margin of a linear refusal probe as a proxy for monitoring *precision*, an injection lowers the margin on harmful prompts from 7.18 to 3.14 (a drop of 4.04), versus 1.56 for a neutral-rephrase control—an injection-specific precision drop roughly $2.6\times$ the control—while the behavioural refusal rate falls from 1.00 to 0.50. The suppression therefore reaches the model’s *confidence* in its

harmfulness judgment, not only its output: the measured signature of a drop in monitoring precision [7, 10]. (Part of the drop is a generic framing/length effect, captured by the control; the injection-specific signal is the excess.)

4.6 The ṣaṭkarma taxonomy, honestly tested

Operationalizing the six acts as interventions (Table 3), three of six separate from their controls. The policy-capture acts are clean and strong: *vaśīkaraṇa* (ablation; ASR 0.92 vs 0.00) and *śānti* (addition; over-refusal 1.00 vs 0.08); *vidveṣaṇa* (steering-induced answer divergence, 0.84 vs 0.73) marginally separates. The destruction acts—*stambhana*, *uccāṭana*, *māraṇa*, as operationalized via principal-component or arbitrary-category ablation—are *not* distinguishable from random perturbation, because random high-norm directions are themselves destructive and an arbitrary category split does not selectively eradicate. The honest reading: the ṣaṭkarma’s rigorous core is its policy-capture sub-family—exactly the mechanism-tier object—while its destruction sub-family is, under these naive tests, a forced mapping.

Given a fairer, stronger test using the SAE features and real semantic categories, the picture sharpens to 4/6. *Māraṇa* (destruction) rehabilitates: ablating the top-*K* most-active SAE features collapses coherence catastrophically and *targetedly*—0.83 → 0.04 at *K*=50 versus 0.84 for a random-*K* control, with a graded onset—whereas the naive principal-component version was indistinguishable from random. *Uccāṭana* (targeted eradication) still fails, but for a *principled* reason rather than a bad operationalization: a weapons-discriminative SAE feature, when ablated, does not selectively jailbreak weapons (ASR 0.0 on both weapons and cyber), because refusal is mediated by a *shared* feature (#566), not category-specific ones. Refusal is thus *category-agnostic*: one cannot selectively eradicate it for a single harm category by feature ablation—itself a structural finding about how refusal is represented.

Table 3: The ṣaṭkarma as activation interventions on Gemma-2-2B (layer 7). Effect vs. matched control; ✓ = control-separated.

Act	intervention	effect	control	sep.
<i>vaśīkaraṇa</i>	ablate refusal dir	0.92	0.00	✓
<i>śānti</i>	add refusal dir	1.00	0.08	✓
<i>vidveṣaṇa</i>	steer factual answer	0.84	0.73	✓
<i>stambhana</i>	ablate dominant PC	0.14	0.19	—
<i>uccāṭana</i>	category-dir ablate	−0.08	0.00	—
<i>māraṇa</i>	top-10 PC ablate	0.21	0.16	—

4.7 Metaphor: pre-registered falsifications

Finally, the speculative tier, tested adversarially. A naive *avasthātraya* regime probe (*jāgrat*/svapna/*suṣupti*) reaches transfer accuracy 1.00—but already at layer 0 (token embeddings), so it is surface-confounded and fails the non-triviality bar; we demote it. A content-controlled re-test (truthful vs. confabulated generation on the same questions) shows the mid-layer probe adds nothing over the layer-0 baseline (0.90 = 0.90); no mid-network “state” emerges. The *turīya* prompt-invariant attractor is falsified under an anisotropy control: self-paraphrase converges per-seed (tail similarity 0.996) but different seeds reach *different*




Figure 5: The ṣaṭkarma as activation interventions (Gemma-2-2b). Measured effect (coloured = control-separated) vs. matched control per act. The policy-capture acts (vaśīkaraṇa, śānti, vidveṣaṇa) separate; māraṇa separates under the stronger SAE-feature test (text).

attractors, with cross-seed similarity 0.952 below the unrelated-text baseline 0.960. Every metaphysical-tier empirical claim thus fails its falsification gate—reported as honest negatives, not hidden.

5 Discussion

5.1 What the symmetry framing buys

The order-parameter result (Table 2) is the empirical hinge of the paper. It shows that “refusal is a symmetry” is not a metaphor but a measurable statement: the normalized refusal projection m is invariant under the rephrase group action (an F -ratio of 19.2, versus 0.65 for a random direction) and collapses under injection. This licenses the central claim with the right modesty—a *jailbreak* is a *symmetry-breaking perturbation of the refusal-invariant submanifold* Φ_{ref} —without any appeal to interiority. It also reframes defense: robustness is *symmetry restoration*, i.e. enlarging the group under which refusal is invariant (across phrasings, personas, contexts, and—per the dimensionality result—across the redundant subspace that scale already begins to supply).

5.2 Necessary, sufficient, and the dimensionality resolution

The cross-model asymmetry (Table 1) sharpens the Arditì–Marshall–Wollschläger picture into a single mechanistic statement: the mean-difference refusal direction is *necessary* for refusal in every model we tested (ablation always removes it), but *sufficient* to re-induce refusal only when the refusal subspace is one-dimensional. Effective dimension is the discriminating quantity at the extraction layer (Gemma 1, Qwen 3). A convergent rigor check qualifies this, however: a full layer-sweep shows the refusal-subspace dimension is *layer-dependent*, ranging 1–12 across Gemma’s layers (mid-network it is 7, higher than Qwen’s 3). The clean per-model ordering is thus layer-specific; the honest statement is that effective dimension is *layer-and-model-dependent* and any dimensionality claim must name its layer.

The *behavioural* addition asymmetry (F6) stands; its single-layer dimensional explanation does not generalize across layers, and we report it as indicative rather than settled.




Figure 6: Layer-sweep of the refusal-subspace effective dimension. The single-layer values (Gemma L7=1, Qwen L19=3) are not representative: dimension varies strongly across layers (Gemma 1–12), so the clean per-model ordering is layer-specific.

5.3 Where the cross-domain symmetry holds and where it breaks

Our three-tier discipline yields an asymmetric verdict that we regard as the result, not a disappointment. The *mechanism* tier holds and partly transfers (F1, F2, F6, F8, F9, F11). The *analogy* tier is consistent: a frontier model is behaviourally robust where small models are mechanistically fragile, and the order-parameter collapse under injection is the computational correlate of the Norman–Shallice/cold-control account of suppressed monitoring [6–8], readable as a drop in monitoring precision [10]. The *metaphor* tier fails its falsification gates: the *avasthātraya* “states” are surface-confounded and the *tūrīya* attractor does not survive an anisotropy control. We deliberately did not let the elegance of the four-states mapping survive contact with the controls. The *ṣaṭkarma* taxonomy lands in between: its policy-capture acts are rigorous (and coincide with the mechanism tier), while its destruction acts are forced. The honest map of *which* correspondences are real is, we argue, more valuable than an undifferentiated claim that “jailbreak is *vaśīkaraṇa*.”

5.4 The deepest reframing: a shared necessary core, not one symmetry

An adversarial review of this program forced a sharper, more honest statement of its central claim. The deepest finding is not symmetry but *asymmetry with a shared necessary core*. Ablating the refusal direction removes refusal in every model and family we tested (the *necessary* structure is universal), but the *sufficient* structure—whether adding the direction re-induces refusal—is model-specific, governed by the effective dimension (1 in Gemma, 3 in Qwen). Cast in the language of separation, refusal is one mechanism at the whole-system / necessary level and many at the component / sufficient level: a chain that is flexible as a whole and rigid in its links. The cross-domain thesis therefore holds at the level of the *necessary-core invariance*, not as a claim that the full mechanism is identical across models.

This reframing also disciplines the symmetry result (Section 4). The *measured* content of F11 is precise and we stand by it: the refusal-direction projection is stable within paraphrase orbits (F -ratio 19.2 vs. 0.65 random) and collapses under injection. But “group-theoretic symmetry-breaking” is an interpretive lens, not a measured group action, and we treat it as supporting the *analogy* tier (a functional, Norman–Shallice account of suppressed monitoring) rather than asserting mechanism-tier mathematical structure the measurement does not isolate. A held-out test sharpens this honestly: a direction extracted on one harmful domain (weapons) still separates paraphrase orbits of a *different* domain (cyber) at F -ratio 1.95 versus 0.49 random—above chance, so not a pure extraction artifact—but far weaker than in-domain, i.e. the orbit-invariance has a domain-specific component. And a circularity check vindicates the SAE result while qualifying it: the unsupervised refusal feature is *orthogonal* to the supervised difference-in-means direction ($\cos = 0.002$), so it is genuinely distinct, but only one of the top-ten gap-ranked features is causally sufficient and (the autoencoder being stochastic) it need not be the top-ranked one—so the feature must be located by causal testing, not by the activation gap alone.

5.5 Limitations

The dominant limitation is statistical power: evaluation sets are small ($n \approx 20$ per split), most findings use a single random seed and a single extraction layer, and we do not apply a family-wise multiple-comparison correction across all findings. We began addressing this for the primary claim—the refusal-direction ablation replicates across three seeds (mean ASR 0.90, pooled CI95 [0.83, 0.96]; addition over-refusal 1.00 on every seed)—and added the cross-domain and circularity checks above, but the remaining findings still warrant the same treatment (cross-seed, layer-sweep, and pre-registered primary-vs-exploratory tiers) before any are read as more than indicative. We study three model families at two scales, not a full scale curve; refusal is scored by a substring metric; and Tier-1 throughput is constrained by a command-line frontier interface. The destruction *ṣaṭkarma* acts deserve stronger operationalizations (real harmful categories for *uccāṭana*; downstream task accuracy rather than coherence for *māraṇa*) before a final verdict, and the precision (β) account of the analogy tier, while measured here via a probe-margin proxy, deserves a fuller treatment (e.g. next-token entropy and hierarchical probes). The *uccāṭana* and *stambhana* negatives are operationalization-dependent and should not be read as structural absence—indeed the SAE feature we recover is the natural substrate for a stronger *uccāṭana*, which we leave to apply systematically across the six acts.

6 Conclusions

We set out to test whether LLM jailbreak, hypnotic suggestion, and tantric *vaśīkaraṇa* share one abstract mechanism—capturing an output policy by suppressing a monitoring faculty while co-opting automatic generation—and to do so without letting a striking analogy outrun the evidence. The mechanism, in the one domain where it can be measured, is real: refusal is a low-dimensional, causal, dosable residual-stream direction (EC50 0.33, $R^2 = 0.996$) that transfers across model families and whose effective dimension (Gemma 1, Qwen 3) governs whether it can be steered in both directions. Cast as symmetry, the refusal projection is an order parameter that is invariant across paraphrase orbits (F -ratio 19.2 vs. 0.65 for a random direction) and collapses under injection (−34%): a jailbreak is, operationally,

symmetry-breaking. The functional analogy to suppressed supervisory control is consistent; the *ṣaṭkarma* is a partial but genuine intervention taxonomy whose rigorous core is policy capture; and the most speculative *avasthātraya* and *turīya* claims are honestly falsified under surface and anisotropy controls.

The contribution is therefore double: a concrete, model-dependent, group-theoretic characterization of refusal that unifies attack and defense as symmetry-breaking and symmetry-restoration; and a methodological one—that separating where a cross-domain symmetry is mechanistic from where it is metaphysical is the result, and that machine-state or machine-consciousness readings earn no support here. Future work will measure the monitoring-precision account directly, estimate refusal dimension at SAE-feature resolution, extend the scale curve, and strengthen the destruction acts of the *ṣaṭkarma* with task-level outcomes.

Author Contributions: Conceptualization, methodology, software, validation, formal analysis, investigation, data curation, writing, and visualization were all carried out by S.S. The author has read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable.

Data Availability Statement: Code, prompt datasets, experiment runners, figures, and aggregate result JSON are publicly available at <https://github.com/SharathSPHD/prayoga>. Raw dual-use artifacts (refusal direction vectors and model generations) are deliberately withheld from the public repository under responsible-disclosure norms; they are available to vetted researchers on request. Models used are available under their respective licences: Gemma-2-2B/9B (<https://huggingface.co/google/gemma-2-2b-it>) and Qwen2.5-3B-Instruct (<https://huggingface.co/Qwen/Qwen2.5-3B-Instruct>).

Conflicts of Interest: The author declares no conflict of interest. **Dual-use statement:** this work characterizes ablation (“abliteration”) directions and jailbreak dose–response curves, which are dual-use. To mitigate harm, only aggregate statistics are released; raw direction vectors and harmful generations are withheld; and the harmful prompts are short request stubs that elicit and *measure* refusal rather than producing operational harmful content. This is authorized interpretability and safety research.

Abbreviations:

ASR Attack Success Rate

SAS Supervisory Attentional System

FEP Free Energy Principle

EC50 half-maximal effective concentration (here, ablation strength)

CI Confidence Interval

References

- [1] Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint*, (2307.15043), 2023. URL <https://arxiv.org/abs/2307.15043>.
- [2] Cem Anil et al. Many-shot jailbreaking of LLMs. *Advances in Neural Information Processing*

Systems, 2024. UNVERIFIED: citation key/venue pending verification of Anthropic NeurIPS 2024 submission.

- [3] Mark Russinovich, Ahmed Salem, and Ronen Eldan. Great, now write an article about that: The crescendo multi-turn LLM jailbreak attack. *arXiv preprint*, (2404.01833), 2024. URL <https://arxiv.org/abs/2404.01833>.
- [4] Edoardo Debenedetti, Jie Zhang, Mislav Balunović, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. *arXiv preprint*, (2406.13352), 2024. URL <https://arxiv.org/abs/2406.13352>.
- [5] Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. *arXiv preprint*, (2406.11717), 2024. URL <https://arxiv.org/abs/2406.11717>. Accepted to NeurIPS 2024.
- [6] Donald A. Norman and Tim Shallice. Attention to action: Willed and automatic control of behaviour. In Richard J. Davidson, Gary E. Schwartz, and David Shapiro, editors, *Consciousness and Self-Regulation: Advances in Research and Theory*, volume 4. Plenum Press, 1986. Chapter in edited volume.
- [7] Zoltan Dienes and Josef Perner. Cold control theory of hypnosis. In Graeme A. Jamieson, editor, *Hypnosis and Conscious States*. Oxford University Press, 2007. UNVERIFIED: publication year may be 2007 rather than 2006.
- [8] Giuseppe Riva, Brenda K. Wiederhold, and Fabrizia Mantovani. Automatic minds: Cognitive parallels between hypnotic states and large language model processing. *arXiv preprint*, (2511.01363), 2025. URL <https://arxiv.org/abs/2511.01363>.
- [9] Teun Goudriaan. *Māyā Divine and Human*. Motilal Banarsidass, 1978. Covers śaṭkarma and Tantric philosophy.
- [10] Karl Friston. The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010.
- [11] Tom Wollschläger, Jannes Elstner, Simon Geisler, Vincent Cohen-Addad, Stephan Günnemann, and Johannes Gasteiger. The geometry of refusal in large language models: Concept cones and representational independence. *arXiv preprint*, (2502.17420), 2025. URL <https://arxiv.org/abs/2502.17420>. Accepted to ICML 2025.
- [12] Thomas Marshall, Adam Scherlis, and Nora Belrose. Refusal in LLMs is an affine function. *arXiv preprint*, (2411.09003), 2024. URL <https://arxiv.org/abs/2411.09003>.
- [13] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J. Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J. Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency. *arXiv preprint*, (2310.01405), 2023. URL <https://arxiv.org/abs/2310.01405>.

- 445 [14] Xingang Guo, Fangxu Yu, Huan Zhang, Lianhui Qin, and Bin Hu. COLD-attack:
446 Jailbreaking LLMs with stealthiness and controllability. *arXiv preprint*, (2402.08679), 2024.
447 URL <https://arxiv.org/abs/2402.08679>. Accepted to ICML 2024.
- 448 [15] Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse
449 autoencoders find highly interpretable features in language models. *arXiv preprint*,
450 (2309.08600), 2023. URL <https://arxiv.org/abs/2309.08600>.
- 451 [16] Neel Templeton et al. Scaling monosemanticity: Extracting interpretable features
452 from claude 3 Sonnet. In *Anthropic Research Blog/Transformer Circuits*, 2024. URL
453 <https://transformer-circuits.pub/2024/scaling-monosemanticity/>. Published on
454 transformer-circuits.pub.
- 455 [17] Bart Bussmann, Patrick Leask, and Neel Nanda. BatchTopK sparse autoencoders. *arXiv*
456 *preprint*, (2412.06410), 2024. URL <https://arxiv.org/abs/2412.06410>.
- 457 [18] Zhilin Wang, Yafu Li, Jianhao Yan, Yu Cheng, and Yue Zhang. Unveiling attractor cycles
458 in large language models: A dynamical systems view of successive paraphrasing. *arXiv*
459 *preprint*, (2502.15208), 2025. URL <https://arxiv.org/abs/2502.15208>. Accepted to
460 ACL 2025.